

Arm ETM/PMU

Haonan Li, Wenxuan Shi, Xueying Zhang

COMPASS

April 20, 2021

① Haonan

② Xueying

③ Wenxuan

1 Plan of Last Week

| 3

- ▶ Syscall capturing: how to handle the content of read
- ▶ Paper writing: revise & rewrite

- ▶ read syscall:
`ssize_t read(int fd, void *buf, size_t count);`
 - > fd: the *file descriptor*
 - > buf with count length: read from fd with count length to buf
 - > return the actual length read to buf
- ▶ original version: only record the return value and the fd
- ▶ record buf [Option 1]: record the whole buf (may incur considerable overhead)
- ▶ record buf [Option 2]: truncate the buf, only first 256 bytes.

```
int main(){
    freopen("record","r",stdin);
    int cnt = 0;
    while(scanf("%s",buf) != -1){
        cnt++;
    }
    printf("%d\n",cnt);
}
```

- ▶ file record: a huge file (2GB), *generated by 100 min of syscall capturing*
- ▶ 493,437 reads
- ▶ typically, the count of read is 4096/8192 bytes (for scanf, 4096 bytes)
though read support 0x7ffff000 bytes

1 Performance Test: Read Huge file

| 6

perform well on an average program (*while there is no read*)

type	real time	record file	Estimated 24-hour file size
baseline	2 min 50.3 s	-	-
syscall capturer (all content)	3 min 11.552 s (+12.5%)	2.0 GB	902.1 GB
syscall capturer (truncate)	2 min 57.675 s (+4.33%)	120 MB	52 GB

1 Performance Test: nginx

| 7

type	real time	record file	Estimated 24-hour file size
baseline	130.758 s	-	-
syscall capturer (all content)	134.193 s (+2.6%)	329 MB	144 GB

Figure 1: Overhead for nginx

1 Plan for Next Week

| 8

- ▶ Continue to revise the paper before delivery to Zhenyu

① Haonan

② Xueying

③ Wenxuan

2 Last Week's plan & This Week's Work

| 10

Write paper

Commits on Apr 18, 2021

background. concurrency bugs. without pic



hotaru555 committed 20 hours ago ✓

background.breakpoint



hotaru555 committed 22 hours ago ✓

Commits on Apr 17, 2021

related work: add 3 , modify 1 examples



hotaru555 committed 2 days ago ✓

related work 2nd attempt



hotaru555 committed 2 days ago ✓

Commits on Apr 14, 2021

2 Next week's plan

| 11

Write paper (add a picture about concurrency bugs and more references).

① Haonan

② Xueying

③ Wenxuan

3 Last week's plan

| 13

- ☐ Writing the non-concurrency evaluation part (still working on)
- ☒ Help writing other parts

224 - supported by `\textit{Control Flow Builder}` and `\textit{Data Flow Builder}` to
 225 - finish the reconstruction of the execution and data of the original application
 226 - in offline analysis phase. Subsequently, based on `\textit{Builder}` results, `\textit{Thwame}`
 227 - uses the root cause detector to diagnose bugs.

228
 229 `\subsubsection{Control Flow Builder}`
 230
 231 - `\textit{Control Flow Builder}` uses the decoded ETW trace result and the binary
 232 - code of the program to reconstruct the actual control flow. The control flow is
 233 - arranged in chronological order by the instructions actually executed by
 234 - different threads of the program. Control flow can help us find the specific
 235 - instruction set that caused the program to crash, and then we can further
 236 - analyze the root cause of the crash. Because the control flow is reconstructed
 237 - based on the trace result, the reconstructed control flow only contains
 238 - instructions in the user space. For multi-threaded programs, we use the
 239 - timestamps in trace result to determine the order of instructions of different
 240 - threads.

241
 242
 243 `\subsubsection{Data Flow Builder}`

244 - `\textit{Data Flow Builder}` restores the data flow corresponding to the control
 245 - flow, and then the actual operation value of each instruction in the control
 246 - flow can be obtained by searching the data flow. Data flow can help us more
 247 - accurately find out the root cause of program crashes. With the data flow, we
 248 - can find out which instructions are operating on the same address, and then find
 249 - the root cause of the crash caused by the concurrent operation.

250 -
 251 - REPT uses the control flow and core dump generated when the program crashed to
 252 - recover the data flow. This method combines reverse deduction and sequential
 253 - deduction to complete the recovery. Because most instructions cannot deduct the
 254 - value before the register or memory is changed, some data in the data flow
 255 - cannot be restored or restored to the wrong value. Although not all data is the
 256 - key information for analyzing bugs, a more accurate data flow can more
 257 - accurately analyze the root cause.

258 -
 259 - Far reducing the possible errors in the data flow, we only use sequential
 260 - derivation to restore the data flow. In order to complete the derivation, the
 261 - initial state of each thread is required. We use breakpoints to output a
 262 - core dump when each thread starts, and assume the content of the core dump to be
 263 - the initial state of the thread.

264 -
 265 - If we can keep the state before the instruction is executed, combined with the
 266 - meaning of the instruction, we can deduce the changes that the instruction will
 267 - cause to the register and memory values. With the initial state of all threads
 268 - and the changes made by each instruction in the control flow, we can restore the
 269 - complete data flow.

270 -
 271 - Because the control flow does not contain instructions executed by the system
 272 - call, we directly use the changes obtained by the syscall capture to help
 273 - restore the data flow.
 274 -

275

276 + reconstruct the control flow and data flow of the original application in the
 277 + offline analysis with the help of `\textit{Control Flow Builder}` and `\textit{Data Flow Builder}`. Subsequently, based on the results, `\textit{Thwame}` uses the root
 278 + cause detector to diagnose bugs.

279
 280 `\subsubsection{Control Flow Builder}`
 281
 282 + `\textit{Control Flow Builder}` uses the decoded ETW trace result and the
 283 + program's binary code to reconstruct the control flow. Control flow represents
 284 + the order of execution instructions, which can help us find the specific
 285 + instruction set that caused the program to crash. Since the control flow is
 286 + reconstructed based on the trace result, it only contains instructions in the
 287 + userspace. For parallel programs, we leverage the timestamps in the trace result
 288 + to determine the order of instructions in different processors.

289
 290
 291 `\subsubsection{Data Flow Builder}`

292 + `\textit{Data Flow Builder}` uses the control flow and core dump, along with a few
 293 + captured online data, to reconstruct the corresponding data flow. The effect of
 294 + each instruction in the control flow can be obtained by searching the data flow.
 295 + Data flow can help us find out the root cause of program crashes more
 296 + accurately. For example, it helps to reveal the root cause of crashes due to
 297 + concurrent operations by locating the instructions operating on the same
 298 + address.

299 +
 300 + comments: Is it reasonable to mention other's work in DESIGN part?
 301 +
 302 + REPT utilizes the control flow and core dump generated by the crashed program to
 303 + recover the data flow `\code{Csuite\csuite\rept}`. This method combines both forward and
 304 + backward execution with error correction to complete the recovery.
 305 + Since most instructions are not reversible (i.e., one cannot trace back the
 306 + registers and memory status), the accuracy of this method cannot be guaranteed.
 307 + % Moreover, not all data contain key information for analyzing bugs. Therefore, a
 308 + % more accurate data flow should be adapted to analyze the root cause.


309 +
 310 + To reduce the possible errors in the data flow, we only use sequential
 311 + derivation to restore the data flow. By combining instructions in the binary
 312 + code and core dump produced at the initial state of each thread, we can evaluate
 313 + the changes to memory and registers for each instruction execution.

314 +
 315 + We use breakpoints to produce a core dump when each thread starts, and assume the
 316 + content of the core dump to be the initial state of the thread. Furthermore, we
 317 + can deduce the changes that the instruction will cause to memory and registers.
 318 + For system calls, due to the impracticality of the kernel instruction trace, we
 319 + restore the system state based on the online capture information.

320 +
 321 + With the initial state of all threads and the changes made by each instruction
 322 + in the control flow, we can restore the complete data flow.

323

- ☐ Abstract
- ☐ Introduction
- ☒ Background
 - ☒ ETM
 - ☐ Concurrency Bug
- ☐ Related Work
- ☐ Design
 - ☒ Online Record (ETM Manager, Syscall Capturer)
 - ☒ Offline Analysis (Control Flow builder, Data Flow builder)
 - ☐ Root Cause Detector
- ☒ Implementation
 - ☐ Online Record (ETM Manager, Syscall Capturer)
 - ☒ Online Record (Library Hook)
 - ☐ Offline Analysis (Control Flow builder, Data Flow builder)
 - ☐ Root Cause Detector
- ☐ Evaluation
- ☐ Conclusion



Don't consider taking the TOEFL/IETLS/GRE test.

3 Next week's plan

| 17

- ☐ Writing the non-concurrency evaluation part (today afternoon)
- ☐ Continue reviewing